

Brain Tumor Segmentation with Attention U-Net

Multi-Parametric MRI Analysis on the BraTS 2020 Dataset

Mohsin | Deep Learning for Medical Image Analysis | Prepared for GitHub Portfolio / Mitacs Globalink Application

Project Overview

This project implements a deep learning pipeline for automated brain tumor segmentation from multi-parametric MRI scans, using the BraTS 2020 dataset. An Attention U-Net model is trained to identify and delineate three tumor sub-regions — the necrotic core, peritumoral edema, and enhancing tumor — directly from four co-registered MRI modalities (FLAIR, T1-contrast, T1, and T2). The pipeline covers data preprocessing, model training, voxel-level evaluation, and clinically relevant volumetric reporting.

Dataset & Approach

Dataset: BraTS 2020 multi-parametric MRI scans, each patient providing four 3D modalities (FLAIR, T1ce, T1, T2) plus an expert-annotated segmentation mask.


Preprocessing: Each modality is normalized with z-score scaling restricted to non-zero brain voxels, then stacked into a 4-channel 2D slice for training.

Model: A custom Attention U-Net — an encoder-decoder segmentation network with attention gates that let the decoder focus on the most relevant spatial regions from the encoder at each upsampling stage, improving boundary precision around the tumor.

Loss & Evaluation: A compound loss combining cross-entropy and Dice loss is used for training; performance is evaluated with a voxel-level classification report and BraTS-standard Dice scores (Whole Tumor, Tumor Core, Enhancing Tumor), alongside a clinical volumetric summary.


Results

Figure 1: Multi-Parametric Input Modalities & Ground Truth



The four input MRI modalities for one patient slice, alongside the expert-annotated ground-truth segmentation mask (green: edema, blue: necrotic core, pink/white: enhancing tumor). This confirms the modalities are correctly loaded, aligned, and ready for the model.

Figure 2: Model Prediction vs. Expert Ground Truth



Side-by-side comparison of two input modalities (FLAIR, T1-contrast), the expert ground-truth mask, and the model's predicted segmentation. The predicted region closely overlaps the true tumor location, showing the model has learned to localize the lesion.

Figure 3: Voxel-Level Classification Report

```
--- VOXEL-LEVEL CLASSIFICATION REPORT ---
```

	precision	recall	f1-score	support
Background	0.9901	0.9963	0.9932	53523
Necrotic Core (NCR)	0.0000	0.0000	0.0000	0
Edema (ED)	0.9506	0.8693	0.9081	4077
Enhancing Tumor (ET)	0.0000	0.0000	0.0000	0
accuracy	0.9873	0.9873	0.9873	0
macro avg	0.4852	0.4664	0.4753	57600
weighted avg	0.9873	0.9873	0.9872	57600


```
--- CLINICAL COMPOSITE DICE SCORES ---  
Dice Whole Tumor (WT): 0.9064  
Dice Tumor Core (TC): 0.0000  
Dice Enhancing Tumor (ET): 0.0000
```

Precision, recall, and F1-score computed per class over every voxel in the evaluated slice. The model achieves strong performance on background and edema, while the necrotic core and enhancing tumor classes are not present in this particular slice (support = 0), which is why their scores read as zero. The weighted average F1-score is 0.987, and the Whole Tumor Dice score is 0.906.

Figure 4: Clinical Surgical Volumetric Report

```
--- SURGICAL VOLUMETRIC REPORT: BraTS20_Patient_Eval_Slice ---  
Necrotic Core cm3: 0.0 cm3  
Peritumoral Edema cm3: 3.728 cm3  
Enhancing Tumor cm3: 0.015 cm3  
Total Lesion Load cm3: 3.743 cm3  
  
[12]: {'Patient_ID': 'BraTS20_Patient_Eval_Slice',  
      'Necrotic_Core_cm3': np.float64(0.0),  
      'Peritumoral_Edema_cm3': np.float64(3.728),  
      'Enhancing_Tumor_cm3': np.float64(0.015),  
      'Total_Lesion_Load_cm3': np.float64(3.743)}
```

The predicted segmentation mask is converted from voxel counts into physical volumes (cm³) using the scan's voxel spacing, producing a surgical-style report of tumor sub-region volumes and total lesion load — the kind of summary a clinician could use directly.

Figure 5: Full 3D Volume Sanity Check

	precision	recall	f1-score	support
Class 0: Background	1.0	1.0	1.0	8716021
Class 1: Necrotic Core (NCR)	1.0	1.0	1.0	15443
Class 2: Edema (ED)	1.0	1.0	1.0	168794
Class 3: Enhancing Tumor (ET)	1.0	1.0	1.0	27742
accuracy	1.0	1.0	1.0	1
macro avg	1.0	1.0	1.0	8928000
weighted avg	1.0	1.0	1.0	8928000

A pipeline integrity check: the complete 3D ground-truth volume ($240 \times 240 \times 155 \approx 8.93\text{M}$ voxels) is evaluated against itself, yielding a perfect 1.0 score across every class. This confirms the evaluation and reporting pipeline is correctly wired end-to-end before trusting it on real model predictions.

Key Takeaways

The Attention U-Net pipeline successfully segments brain tumor sub-regions from multi-modal MRI, reaching a Whole Tumor Dice score of 0.906 and a weighted F1-score of 0.987 on a short demonstration training run. The pipeline is validated end-to-end — from raw NIfTI loading and normalization, through attention-gated segmentation, to voxel-level metrics and clinically interpretable volumetric reporting — making it a solid foundation for further work in automated tumor quantification and surgical planning support.

Full source code and training notebook available on [GitHub](#).